

Diffusion maps, spectral clustering and reaction coordinates of dynamical systems

Boaz Nadler
Stéphane Lafon
Ronald R. Coifman
Ioannis G. Kevrekidis

Shane McQuarrie

CSE 392: SCIENTIF COMP MACH/DEEP LRN
15 November 2019



The University of Texas at Austin
Oden Institute for Computational
Engineering and Sciences

Outline

- 1 Overview
- 2 Markov Chains, Random Walks, and Diffusion
- 3 Diffusion Maps
- 4 Numerical Examples



Available online at www.sciencedirect.com



Appl. Comput. Harmon. Anal. 21 (2006) 113–127

Applied and
Computational
Harmonic Analysis

www.elsevier.com/locate/acha

Diffusion maps, spectral clustering and reaction coordinates of dynamical systems

Boaz Nadler^{a,*}, Stéphane Lafon^{a,1}, Ronald R. Coifman^a, Ioannis G. Kevrekidis^b

^a Department of Mathematics, Yale University, New Haven, CT 06520, USA

^b Chemical Engineering and PACM, Princeton University, Princeton, NJ 08544, USA

Received 28 October 2004; revised 10 February 2005; accepted 29 July 2005

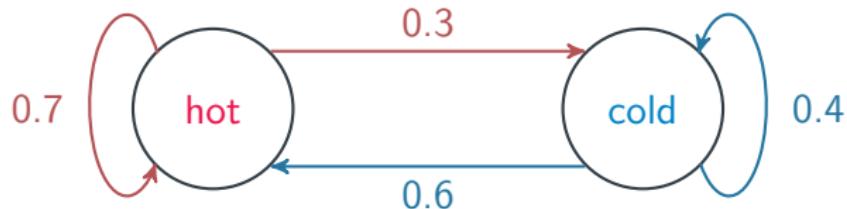
Available online 9 June 2006

Communicated by the Editors

Overview of Paper

- **Statistics:** published in 2006, ~ 500 citations (companion to a paper with $\sim 2,100$ citations)
- **Problem statement:** high-dimensional data coming from large-scale dynamical systems often have low-dimensional structure, but learning this structure requires identifying slow variables and dynamically meaningful reaction coordinates.
- **Solution:** define a family of *diffusion maps* that embed the data in a low-dimensional space with a desired time scale.

Markov Chains



$$\begin{matrix} & \text{hot tomorrow} & \text{cold tomorrow} \\ \text{hot today} & \left[\begin{array}{cc} 0.7 & 0.3 \\ 0.6 & 0.4 \end{array} \right] & = M \\ \text{cold today} & & \end{matrix}$$

Markov Chains

- Let $\mathbf{x} \in \mathbb{R}^n$ be a state distribution vector, i.e., x_j is the probability of being in state j at the current time.
- $(M^\top \mathbf{x})_j$ is the probability of being in state j after one transition
- $((M^t)^\top \mathbf{x})_j$ is the probability of being in state j after t transitions.

Key observation: M^t describes the dynamics with the time scale t .

Key property: M has a steady-state distribution $\psi \succ 0$, i.e., $M^\top \psi = \psi$.

Markov Chains

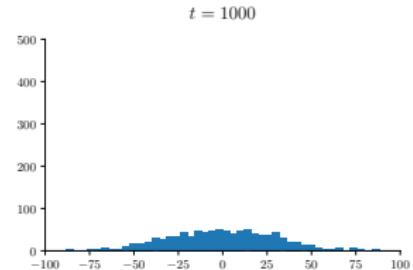
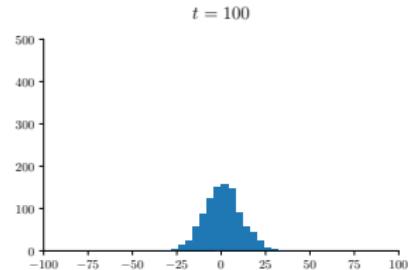
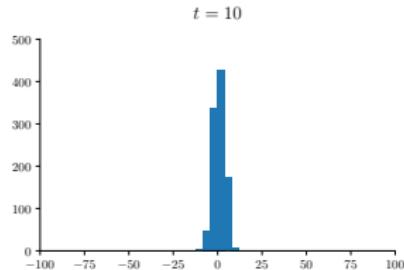
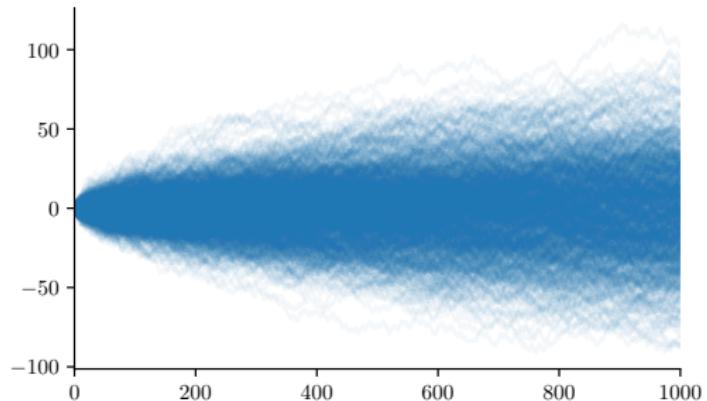
- Discrete case: P_{ij} is the probability of moving from state i to state j .
Therefore,

$$\sum_{j=1}^N P_{ij} = 1 \quad \forall j = 1, 2, \dots, N.$$

- Continuous case: $p(\mathbf{x}, \mathbf{y})$ is a probability density function describing the transition from \mathbf{x} to \mathbf{y} , with

$$\int_X p(\mathbf{x}, \mathbf{y}) d\mu(\mathbf{y}) = 1 \quad \forall \mathbf{x} \in X.$$

Random Walks and Diffusion



From Data to Markov Chain

Let $X = \{\mathbf{x}\}_{i=1}^N \subset \mathbb{R}^n$ be the data, and let $k : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ be a *kernel function* satisfying

$$k(\mathbf{x}, \mathbf{y}) = k(\mathbf{y}, \mathbf{x})$$

$$k(\mathbf{x}, \mathbf{y}) \geq 0$$

$$k(\mathbf{x}, \mathbf{x}) = 0 \quad (\text{usually})$$

Example: Gaussian kernel, $k_\varepsilon(\mathbf{x}, \mathbf{y}) = e^{-\frac{\|\mathbf{x}-\mathbf{y}\|^2}{\varepsilon}}$

Define $M \in \mathbb{R}^{N \times N}$ by

$$M_{ij} = \frac{k(\mathbf{x}_i, \mathbf{x}_j)}{\sum_{\ell=1}^N k(\mathbf{x}_i, \mathbf{x}_\ell)}.$$

Then M_{ij} is the weighted distance from \mathbf{x}_i to \mathbf{x}_j .

Diffusion on Discrete Markov Chains

<https://tinyurl.com/y5u3jhte>

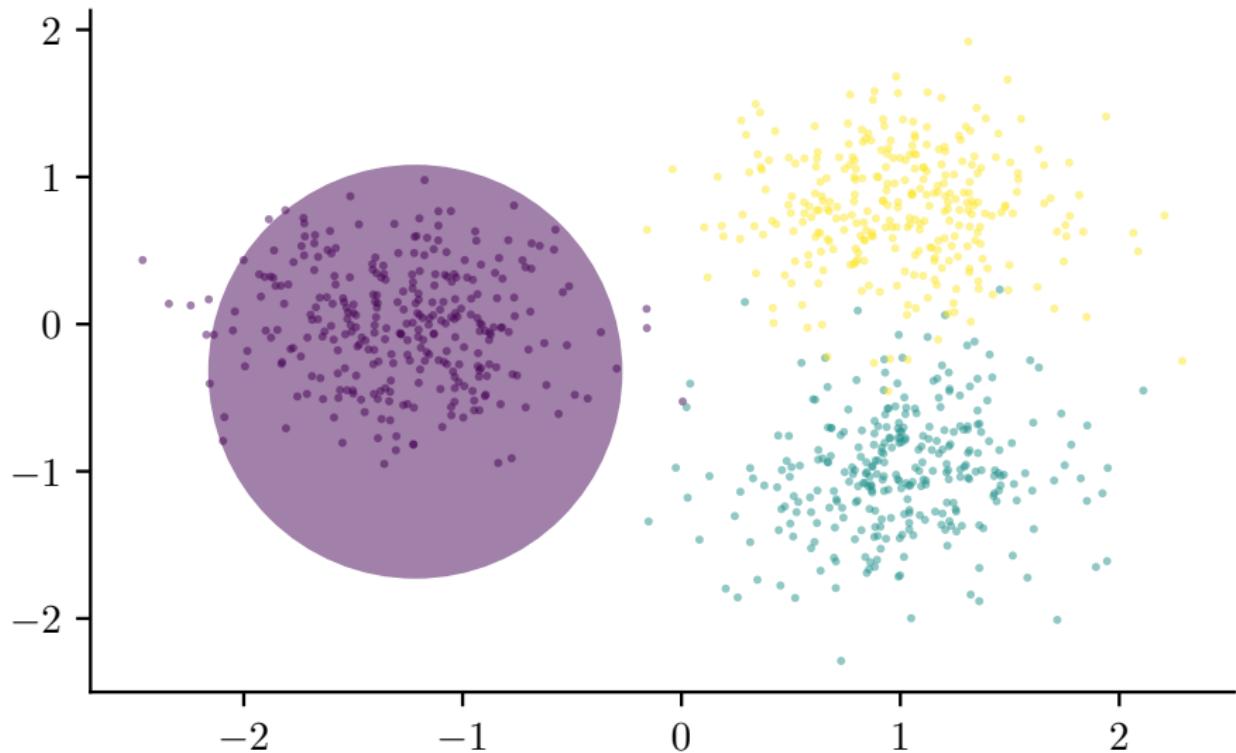
```
import numpy as np

def gaussian_kernel(x, y, eps=.7):
    return np.exp(-np.sum((x - y)**2) / eps)

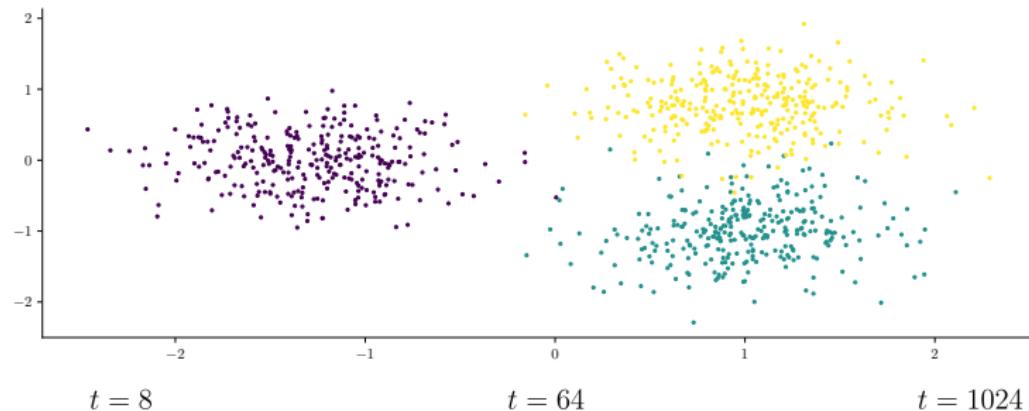
def markov_transition_matrix(X, k, **kwargs):
    m,n = X.shape
    G = np.zeros((m,m))
    for i in range(m):
        for j in range(i+1,m):
            G[i,j] = k(X[i], X[j], **kwargs)
    G = G + G.T
    return G / G.sum(axis=1).reshape((-1,1))
```

Diffusion on Discrete Markov Chains

$$t = 0$$



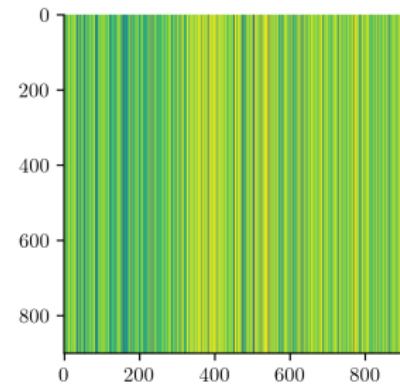
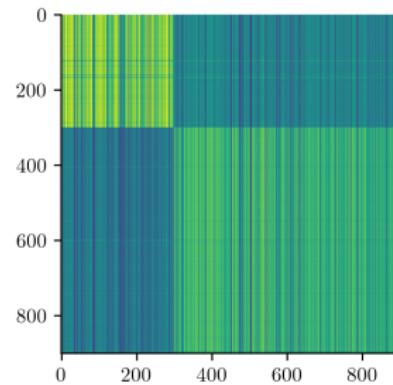
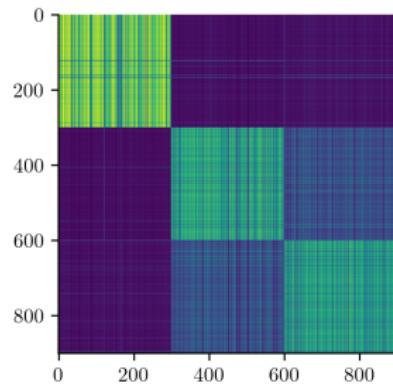
Diffusion on Discrete Markov Chains



$t = 8$

$t = 64$

$t = 1024$



Diffusion Maps

Definition (Diffusion Distance)

Let M be the row-stochastic Markov transition matrix defined previously. Let $\{\lambda_j\}_{j=0}^{N-1}$ be the nondecreasing eigenvalues of M , with $\lambda_0 = 1$, and let $\{\psi_j\}_{j=0}^{N-1}$ be the corresponding right eigenvectors. The *diffusion distance* at time t is given by

$$\begin{aligned} D_t^2(\mathbf{x}, \mathbf{y}) &:= \|p(\mathbf{z}, t|\mathbf{x}) - p(\mathbf{z}, t|\mathbf{y})\|_w^2 \\ &= \sum_{\mathbf{z}} (p(\mathbf{z}, t|\mathbf{x}) - p(\mathbf{z}, t|\mathbf{y}))^2 w(\mathbf{z}). \end{aligned}$$

Choosing $w(\mathbf{z}) = 1/p(\mathbf{z})$, we also have

$$D_t^2(\mathbf{x}, \mathbf{y}) = \sum_j \lambda_j^{2t} (\psi_j(\mathbf{x}) - \psi_j(\mathbf{y}))^2.$$

Wait, what?

Bad Notation

If the data set is $\{\mathbf{x}_i\}_{i=1^N}$, then $\psi_j(\mathbf{x}_i) = (\psi_j)_i$.

(We can't apply this mapping to arbitrary data.)

Diffusion Maps

Definition (Diffusion Map)

Let $\{(\lambda_j, \psi_j)\}_{j=0}^{N-1}$, be the right eigenpairs of M as before. For some $k < N$, we define the family of *diffusion maps* $\Psi_t : X \rightarrow \mathbb{R}^{k+1}$ parametrized by t :

$$\Psi_t(\mathbf{x}) = \begin{bmatrix} \lambda_0^t \psi_0(\mathbf{x}) \\ \lambda_1^t \psi_1(\mathbf{x}) \\ \vdots \\ \lambda_k^t \psi_k(\mathbf{x}) \end{bmatrix}, \quad \text{i.e.,} \quad \Psi_t(\mathbf{x}_i) = \begin{bmatrix} \lambda_0^t (\psi_0)_i \\ \lambda_1^t (\psi_1)_i \\ \vdots \\ \lambda_k^t (\psi_k)_i \end{bmatrix}.$$

Note that $\|\Psi_t(\mathbf{x}) - \Psi_t(\mathbf{y})\| = D_t(\mathbf{x}, \mathbf{y})$ (the diffusion distance).

Diffusion Maps

```
from scipy import linalg as la

class DiffusionMap:
    def __init__(self, X, k, kernel=gaussian_kernel, **kwargs):
        self.M = markov_transition_matrix(X, kernel, **kwargs)
        vals, vecs = la.eig(self.M)
        index = vals.real.argsort()[:-1][:-s]
        self.vals = np.minimum(vals[index].real, 1)
        self.vecs = vecs[:,index].real

    def __call__(self, t):
        return self.vals**t * self.vecs
```

Special Cases

- $\mathbf{x}_i \sim p(\mathbf{x})$ with $p(\mathbf{x}) = e^{-U(\mathbf{x})}$ for some potential $U(\mathbf{x})$.
- \mathbf{x}_i is a state of the stochastic equation $\dot{\mathbf{x}} = -\nabla U(\mathbf{x}) + \sqrt{2}\dot{\mathbf{w}}$.

Main result: Depending on how the transition matrix M is normalized, the eigenvectors of M approach the eigenfunctions of an appropriate diffusion operator as the number of samples increases.
(But this doesn't matter for implementation.)

Summary

Given data $X = \{\mathbf{x}_i\}_{i=1}^N \subset \mathbb{R}^n$ and a kernel $k : X \times X \rightarrow \mathbb{R}$,

- ① Calculate the Markov transition matrix M with

$$M_{ij} = \frac{k(\mathbf{x}_i, \mathbf{x}_j)}{\sum_{\ell=1}^N k(\mathbf{x}_i, \mathbf{x}_\ell)}.$$

- ② Compute the eigenvalue-eigenvector pairs $\{(\lambda_j, \psi_j)\}_{j=0}^{N-1}$ of M
- ③ Define the diffusion map (parametrized by time)

$$\Psi_t(\mathbf{x}_i) = [\begin{array}{cccc} \lambda_0^t(\psi_0)_i & \lambda_1^t(\psi_1)_i & \cdots & \lambda_k^t(\psi_k)_i \end{array}]^\top.$$

Cost: $\mathcal{O}(N^2) + \mathcal{O}(kN^2) + \mathcal{O}(Nk) = \mathcal{O}(kN^2)$
(cost of computing k eigenpairs).

Judgements

Pros

- Correlations between all data points are treated simultaneously
- Gracefully exposes the low-dimensional geometry of a dataset
- Provides interpretation for various timescales
- Relatively cheap (eigenvalue problem)
- Robust to noise

Cons

- Diffusion map only defined on X , not \mathbb{R}^n , so it's hard to deal with new points; projection is possible but hard
- Sometimes the parametrization by time doesn't pay off (but this can also be informative)
- Sometimes not quite so robust to noise (?)

Numerical Examples

<https://tinyurl.com/y5u3jhte>

References

-  Ronald R Coifman and Stéphane Lafon.
Diffusion maps.
Applied and computational harmonic analysis, 21(1):5–30, 2006.
-  Diffusion map.
https://en.wikipedia.org/wiki/Diffusion_map.
Accessed: 2019-11-08.
-  Boaz Nadler, Stéphane Lafon, Ronald R Coifman, and Ioannis G Kevrekidis.
Diffusion maps, spectral clustering and reaction coordinates of dynamical systems.
Applied and Computational Harmonic Analysis, 21(1):113–127, 2006.